

Assignment - 4

Ans-1) A shared Family Notebook.

Imagine The Race Condition:

- 1) Person A checks list and sees "milk, eggs"
- 2) At the same time, Person B checks list and also sees "milk, eggs".
- 3) Person A adds "bread" writes down "milk, eggs, bread."
- 4) Person B adds "butter" and, working from their original mental copy, overwrites the list with "milk, eggs, butter".
- 5) Result: The item "bread" is lost. The final list is wrong because both people were updating the shared resource w/o coordination.

How Mutual Exclusion addresses it.

The solution: A lock on the notebook.

A lock is placed on the notebook. The rule is! you must have the lock to read or write the list.

<u>Ans-3) Feature</u>	Peterson's Solution	Semaphores
Implementation complexity	High. The programmer must manually and correctly code the entire algorithm for each critical section.	Low. The OS provides ready-to-use wait() and signal() operators.
Hardware dependency	High. Relies on atomic hardware instructions for correctness and is vulnerable to modern CPU memory reordering.	Low (for dry user). The hardware dependent component is buffer wins.

Ans-3) One advantage of using monitors is that they integrate the lock and condition variable, making the code less error-prone.

- with semaphores, a programmer can easily cause deadlock by incorrectly ordering the wait operations.
- with monitors, the lock is automatically acquired on entry and released during condition waits. This atomic "check-and-wait" operating prevents the deadlock seen from in semaphore-based solution, leading to more robust code in multi-core systems where concurrent access is the norm.

Ans 4) How starvation occurs

In the common "Readers-preference" solution, a continuous stream of new readers can enter the critical section as long as at least one reader is inside. A waiting writer must wait for all readers to finish, which may never happen if new readers keep arriving.

One method to Prevent it.

Implement a "fair" sell using a turnstile (like a semaphore).

- Both readers and writer must wait on this turnstile semaphore.
- A writer who acquires it will block all new readers from starting until the writer has finished.
- This ensures a first-come, first-served order, preventing a stream of readers from indefinitely slowing a writer.

Ans-5) Practical Drawback: Severe Reduction in Resource utilization

The common method to eliminate "Hold and wait" is to request a process to request all needed resources at once before it starts.

The drawback is that resources allocated

to a process may sit idle for longer period. For example, a printer first needs a scanner now and a plotter later must lock both at the start. The plotter is unavailable to other processes while the process uses the scanner, leading to poor overall system performance and throughput.

Ans-b) a) The global graph is formed by connecting the local fragments

- From S1: $P_1 \rightarrow P_2$ and $P_3 \rightarrow P_4$
- From S2: $P_2 \rightarrow P_5$ and $P_5 \rightarrow P_6$
- From S3: $P_6 \rightarrow P_1$

$$P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6 \rightarrow P_1$$

b) Yes, a deadlock exists. The processes involved in the circular wait are P_1, P_2, P_5 and P_6 .

c) A suitable algorithm is the Path-Pushing algorithm. In this method, a site that suspects a deadlock initiates a probe message. This probe travels along the edges of the wait-for graph. If the probe ever returns to the

Ans)
$$\begin{aligned} \text{Expected Time} &= (\text{Probability local} \times \text{Time local}) \\ &\quad + (\text{Probability remote} \times \text{Time Remote}) \\ &= (0.7 \times 5\text{ms}) + (0.3 \times 25\text{ms}) = 3.5 + 7.5 = 10\text{ms}. \end{aligned}$$

b) Strategy: Client-side caching

Justification: After a file is accessed remotely once, subsequent reads are served from the local cache. This reduces the effective probability of a remote access, dramatically lowering the average access time from 25 ms to 5 ms for cached files.

Ans-9) a) Total overhead : $200ms + (9 \times 50ms) = 650ms$

b) The 1-second RPO requires a checkpoint every 1 second using one full checkpoint as a base and nine faster incrementals. is the cheapest valid way to create this continuous recovery char

A Ans-1) a) 1 full + 9 incremental checkpoints
overhead = $200 + (9 \times 50) = 650ms$.

- b) 1-second RPO requires a checkpoint every second
- 1 full checkpoint creates a valid recovery base
 - 9 incremental checkpoints marks the RPO at minimum cost
 - This is the cheapest valid strategy meeting the RPO requirement.